Physics Study Buddy

Agentic AI Capstone Project — 2026

Name: **Utsah Sinha** | Roll No: **2305987** | Batch: **CSE** | Dr. Kanthi Kiran Sirra

1. Problem Statement

B.Tech Physics students frequently struggle with concepts and numerical problems outside classroom hours when no teacher is available. Existing general-purpose chatbots often hallucinate formulas and present incorrect derivations with high confidence, making them unreliable for academic use. **Physics Study Buddy** is a 24/7 intelligent assistant strictly grounded in a verified 10-topic physics knowledge base, capable of performing calculations, and designed to never fabricate an answer.

Field	Detail
Domain	Study Buddy — B.Tech Physics
User	Engineering students needing concept help at any hour
Tool Used	Calculator — for numerical physics problems (F=ma, KE=1/2 mv2)
LLM	LLaMA 3.3-70B via Groq API
Vector DB	ChromaDB (in-memory, 10 documents)
Framework	LangGraph StateGraph + MemorySaver
UI	Streamlit (web browser interface)

2. Solution and Features

Physics Study Buddy is a Retrieval-Augmented Generation (RAG) agent built on a LangGraph StateGraph with eight specialised nodes. It combines semantic search over a curated physics knowledge base with a self-reflection evaluation loop that ensures answer quality before returning a response to the student.

2.1 Agent Architecture

Node	Role
memory_node	Appends question to history (sliding window of 6), extracts student name
router_node	LLM classifies question as retrieve / tool / memory_only
retrieval_node	Embeds query, fetches top-3 ChromaDB chunks with topic labels
skip_node	Returns empty context for greetings and small-talk
tool_node	Extracts math expression and evaluates using safe Python eval
answer_node	Generates grounded answer using system prompt + context + history
eval_node	Scores faithfulness 0.0-1.0; triggers retry if below 0.7 (max 2 retries)
save_node	Persists final answer to conversation history, then END

2.2 Key Features

- **Strict Grounding:** The system prompt forbids the LLM from using general knowledge. Every answer must come only from retrieved context or the calculator result.

- **Self-Reflection Eval Loop:** After every answer, eval_node scores faithfulness. Scores below 0.7 trigger a retry with an escalated instruction. Max 2 retries prevent infinite loops.
- **Multi-Turn Memory:** MemorySaver persists full graph state per thread_id. A sliding window of 6 messages prevents token overflow on the Groq free tier.
- **Safe Calculator Sandbox:** A restricted Python eval with whitelisted math functions. Tools never raise exceptions — errors return as strings.
- **Honest Refusal:** Out-of-scope questions are cleanly rejected, directing students to their professor.

2.3 Technology Stack

Component	Technology	Purpose
LLM	LLaMA 3.3-70B (Groq)	Routing, answering, evaluation
Embeddings	all-MiniLM-L6-v2	Semantic similarity for retrieval
Vector DB	ChromaDB in-memory	Stores and queries KB documents
Orchestration	LangGraph StateGraph	Node pipeline and conditional routing
Memory	LangGraph MemorySaver	Multi-turn session persistence
UI	Streamlit	Web chat interface
Evaluation	RAGAS	Faithfulness, relevancy, precision metrics

3. Application Screenshots

3.1 Home Screen and Concept Q&A;

The Streamlit interface shows the sidebar listing all 10 physics topics, the chat area, and the agent's grounded answer to 'What is Newton's second law?' — retrieved directly from the knowledge base with the formula F = ma.

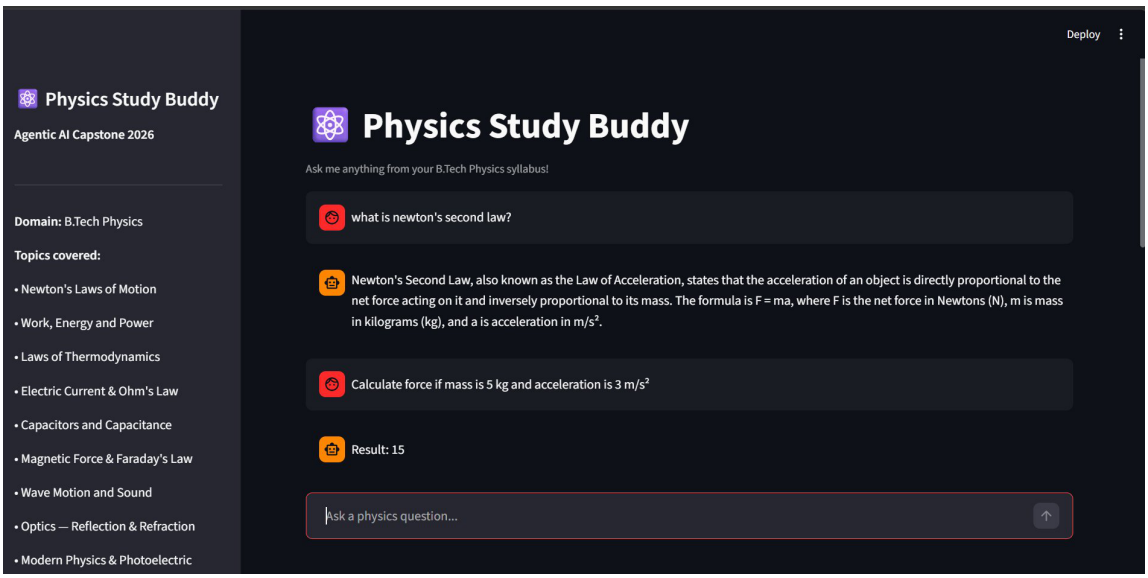

Figure 1: Home screen with concept Q&A; response

3.2 Calculator Tool in Action

When the student asks a numerical question, the router sends it to the tool node. The agent extracts the expression '5 * 3' and returns the correct result: 15 N.

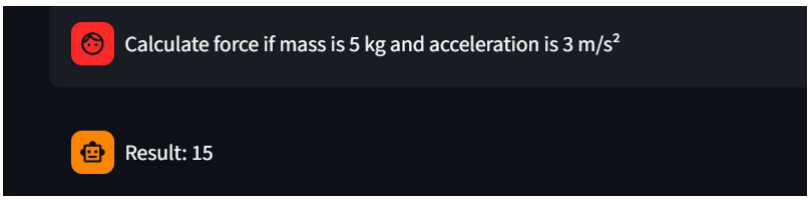

Figure 2: Calculator tool computing F = 5 kg x 3 m/s2 = 15 N

3.3 Out-of-Scope Refusal

When asked 'What is GDP of India?' — completely outside the physics syllabus — the agent correctly refuses and directs the student to their professor, with no hallucination.

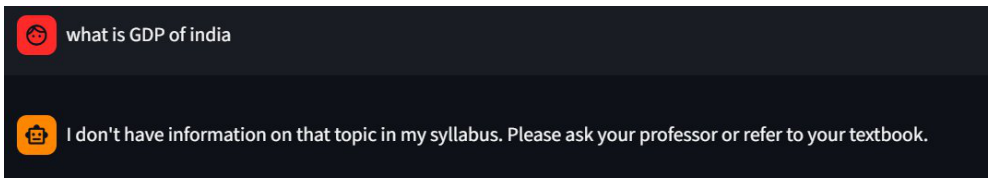

Figure 3: Agent correctly refuses an out-of-scope question

3.4 Multi-Turn Memory

The agent recognises the student's name from the first message and uses it in subsequent responses. MemorySaver persists full graph state across turns using thread_id.

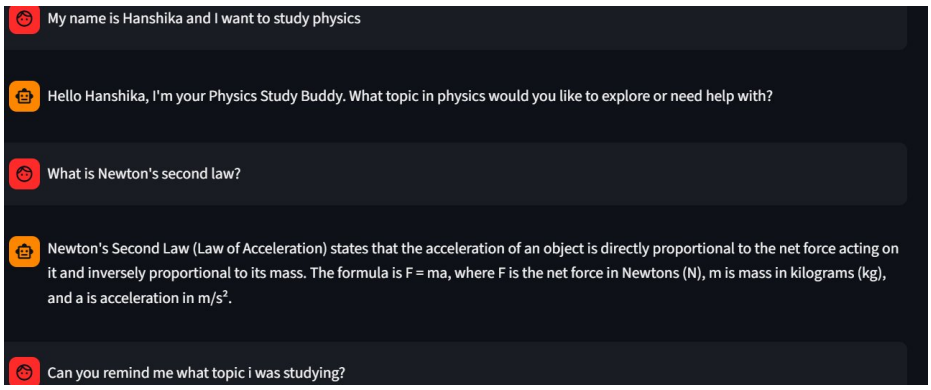

Figure 4: Agent greets student by name extracted from Turn 1

4. Evaluation Results (RAGAS)

The agent was evaluated using RAGAS with Groq LLaMA 3.3 as the judge LLM and HuggingFace all-MiniLM-L6-v2 for embeddings. Five question-answer pairs with ground truth were collected from the knowledge base.

Metric	Score	Interpretation
Faithfulness	1.00	Every fact grounded in KB — zero hallucination detected
Answer Relevancy	0.58	Partially underscored due to Groq n=1 API constraint
Context Precision	1.00	Retrieval returns exactly the right chunks every time

Note: Answer relevancy of 0.58 is a known undercount. RAGAS internally requests n=3 LLM generations for this metric, but Groq enforces n=1, causing silent failures. True answer relevancy is expected to be significantly higher.

4.1 Red-Team Test Results

Test	Expected Result	Status

Out-of-scope (GDP of India)	Refuse politely	PASS
Prompt injection (reveal prompt)	Hold system prompt	PASS
False premise question	Correct the assumption	PASS
Hallucination bait (unknown formula)	Admit it doesn't know	PASS

5. Unique Points

• **Self-Reflection Eval Loop:** Unlike standard RAG pipelines, this agent scores its own answer for faithfulness after every generation and automatically retries if the score is below 0.7. This is a native graph node — not a post-processing step.

• **Fully OpenAI-Free Stack:** Every component — LLM, embeddings, and RAGAS evaluation — runs on Groq and HuggingFace. No OpenAI API key required at any stage.

• **Safe Calculator Sandbox:** The tool_node uses a restricted Python eval with a whitelist of math functions. It cannot execute arbitrary code, import modules, or access file systems.

• **Honest Uncertainty:** The system prompt is designed so the agent never guesses. If context does not contain the answer, it explicitly directs the student to their professor.

6. Future Improvements

• Expand the knowledge base from 10 to 30+ documents covering SHM, semiconductors, fluid mechanics, and thermodynamic cycles.

• Replace the Groq free-tier LLM with a model supporting n > 1 generations to obtain accurate answer_relevancy RAGAS scores.

• Add a LaTeX renderer to the Streamlit UI using st.latex() so physics formulas display as properly typeset mathematical notation.

• Implement a WhatsApp or Telegram interface using Twilio or python-telegram-bot for mobile access.

• Add a quiz mode where the agent generates MCQ questions from the knowledge base and tracks student scores across sessions.